

# Linear Regression Cost Function

## 1. Hypothesis/Output Function

In linear regression, the hypothesis is assumed to be a linear function of the input:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

In vectorized form:

$$h_{\theta}(x) = X \cdot \theta$$

where:

- $X$  is the feature matrix of size  $m \times (n + 1)$ ,
- $\theta$  is the parameter vector of size  $(n + 1) \times 1$ ,
- $h_{\theta}(x)$  is the predicted value for the input  $x$ .

## Hypothesis for Linear Regression

For a single training example  $x^{(i)}$ , the hypothesis function is:

$$h_{\theta}(x^{(i)}) = \theta_0 + \theta_1 x_1^{(i)} + \theta_2 x_2^{(i)} + \cdots + \theta_n x_n^{(i)}$$

This can also be written in vector form as:

$$h_{\theta}(x^{(i)}) = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}^T \cdot \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} = \theta^T \cdot x^{(i)}$$

Here:

- $x^{(i)} = [x_0^{(i)}, x_1^{(i)}, \dots, x_n^{(i)}]^T$  is the feature vector for the  $i$ -th training example (with  $x_0 = 1$  for the bias term).
- $\theta = [\theta_0, \theta_1, \dots, \theta_n]^T$  is the parameter vector.

The result  $h_{\theta}(x^{(i)})$  is a scalar value that represents the predicted output for the  $i$ -th training example.

## Hypothesis for All Training Examples

For  $m$  training examples, we stack the feature vectors  $x^{(1)}, x^{(2)}, \dots, x^{(m)}$  into a single matrix  $X$ , called the **feature matrix**:

$$X = \begin{bmatrix} x_0^{(1)} & x_1^{(1)} & \dots & x_n^{(1)} \\ x_0^{(2)} & x_1^{(2)} & \dots & x_n^{(2)} \\ \vdots & \vdots & \ddots & \vdots \\ x_0^{(m)} & x_1^{(m)} & \dots & x_n^{(m)} \end{bmatrix}$$

Here:

- $X$  is an  $m \times (n + 1)$  matrix:
  - $m$ : Number of training examples (rows).
  - $n + 1$ : Number of features, including the bias term (columns).
- Each row of  $X$  corresponds to a single training example  $x^{(i)}$ .

For all training examples, the predictions can be computed as:

$$\hat{y} = X \cdot \theta$$

where:

- $X$  is the feature matrix of size  $m \times (n + 1)$ .
- $\theta$  is the parameter vector of size  $(n + 1) \times 1$ .
- $\hat{y}$  is the prediction vector of size  $m \times 1$ , where each element is:

$$\hat{y}_i = h_{\theta}(x^{(i)}) = \sum_{j=0}^n \theta_j x_j^{(i)}$$

## 2. Error Between Prediction and Actual Value

For each training example  $i$ , the error is defined as the difference between the predicted value  $h_{\theta}(x^{(i)})$  and the actual target value  $y^{(i)}$ :

$$\text{Error}_i = h_{\theta}(x^{(i)}) - y^{(i)}$$

## 3. Measuring the Error

To evaluate the performance of the model, we use the sum of squared errors (SSE):

$$\text{SSE} = \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

Squaring the errors ensures that:

- Positive and negative errors do not cancel each other out.
- Larger errors are penalized more heavily than smaller errors.

## 4. Averaging the Error

To make the error independent of the number of training examples  $m$ , we calculate the mean squared error (MSE):

$$\text{MSE} = \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

## 5. Cost Function Definition

The cost function  $J(\theta)$  is defined as half the Mean Squared Error (MSE):

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

The factor of  $\frac{1}{2}$  simplifies the gradient computation during optimization, as it cancels the factor of 2 that arises from differentiation.

## 6. Vectorized Form

In vectorized notation, the cost function can be written as:

$$J(\theta) = \frac{1}{2m} \left( (X \cdot \theta - y)^T \cdot (X \cdot \theta - y) \right)$$

where:

- $X \cdot \theta$  computes the predicted values for all training examples,
- $X \cdot \theta - y$  is the vector of residuals (differences between predictions and actual values),
- $(X \cdot \theta - y)^T \cdot (X \cdot \theta - y)$  computes the sum of squared residuals.

## 7. Final Cost Function

The final cost function is:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

## Gradient Descent for Linear Regression

Gradient descent is an iterative optimization algorithm used to minimize the cost function  $J(\theta)$  in linear regression. The updates to the parameter vector  $\theta$  are performed by moving in the direction of the negative gradient of  $J(\theta)$ .

### Cost Function

The cost function for linear regression is defined as:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

### Gradient of the Cost Function

To minimize  $J(\theta)$ , we compute the gradient of  $J(\theta)$  with respect to  $\theta$ :

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)}$$

In vectorized form, this gradient can be written as:

$$\nabla_{\theta} J(\theta) = \frac{1}{m} X^T (X\theta - y)$$

where:

- $X$ : Feature matrix of size  $m \times (n + 1)$ .
- $\theta$ : Parameter vector of size  $(n + 1) \times 1$ .
- $y$ : Output vector of size  $m \times 1$ .
- $X\theta - y$ : The vector of residuals (differences between predicted and actual values).

### Gradient Descent Update Rule

Using the gradient, we update  $\theta$  iteratively:

$$\theta := \theta - \alpha \nabla_{\theta} J(\theta)$$

where:

- $\alpha$ : Learning rate, which determines the step size.

Substituting the gradient:

$$\theta := \theta - \frac{\alpha}{m} X^T (X\theta - y)$$

### Step-by-Step Computation (Vectorized Form)

1. Compute the predictions:

$$\hat{y} = X\theta$$

2. Compute the residuals:

$$\text{residuals} = X\theta - y$$

3. Compute the gradient:

$$\nabla_{\theta} J(\theta) = \frac{1}{m} X^T (X\theta - y)$$

4. Update the parameters:

$$\theta := \theta - \alpha \nabla_{\theta} J(\theta)$$